

# CS 466/566

# Introduction to Deep Learning

Lecture 2

Introduction to Machine Learning - Part 2

Composed from various sources such as Andrew Ng, Quoc Le, etc.

# Logistic Regression (Binary Classification)

- In linear regression we tried to predict the value of  $y^{(i)}$  for the  $i^{\text{th}}$  example  $x^{(i)}$  using a linear function  $y = h_{\theta,b}(x) = \theta^T x + b$
- This is clearly not a great solution for predicting binary-valued labels such as  $y^{(i)} \in \{0, 1\}$ .
- In logistic regression we use a different hypothesis class to try to predict the probability that a given example belongs to the “1” class versus the probability that it belongs to the “0” class.

# Logistic Regression

- Specifically, we will try to learn a function of the form:

$$P(y = 1|x) = h_{\theta,b}(x) = \frac{1}{1 + \exp(-\theta^T x - b)} \equiv \sigma(\theta^T x + b)$$

$$P(y = 0|x) = 1 - P(y = 1|x) = 1 - h_{\theta,b}(x)$$

- $\sigma(z) = \frac{1}{1+e^{-z}}$  is called the “sigmoid” or “logistic” function.
- it is an S-shaped function that “squashes” the value of  $\theta^T x + b$  into the range  $[0, 1]$  so that we may interpret  $h_{\theta,b}(x)$  as a probability.

# Logistic Regression

- Our goal is to search for a value of  $\theta$  so that the probability  $P(y = 1|x) = h_{\theta,b}(x)$  is large when  $\mathbf{x}$  belongs to the “1” class and small when  $\mathbf{x}$  belongs to the “0” class.
- For a set of training examples with binary labels  $\{x^{(i)}, y^{(i)} : i = 1, \dots, m\}$  the following cost function measures how well a given  $h_{\theta,b}$  does this:

$$J(\theta, b) = - \sum_i \left( y^{(i)} \log \left( h_{\theta,b}(x^{(i)}) \right) + (1 - y^{(i)}) \log \left( 1 - h_{\theta,b}(x^{(i)}) \right) \right)$$

# Logistic Regression

$$J(\theta, b) = - \sum_i \left[ y^{(i)} \log \left( h_{\theta, b}(x^{(i)}) \right) + (1 - y^{(i)}) \log \left( 1 - h_{\theta, b}(x^{(i)}) \right) \right]$$

- only one of the two terms in the summation is non-zero for each training example.
- we now have a cost function that measures how well a given hypothesis  $h_{\theta}$  fits our training data.
- we can learn to classify our training data by minimizing  $J(\theta, b)$  to find the best choice of  $\theta, b$ .
- we can classify a new test point as “1” or “0” by checking which of these two class labels is most probable.

# A case study of movie recommendations


- Should I watch Gravity or not?
- We ask our close friends, let's say, Mary and John. They watched it.
- In a scale of 1 to 5, they both rate it as a 3.
- 3 meaning that “not outstanding but worth watching”.
- How can I decide if I should go to Gravity or not?
- I need more data!

# More movie data from friends

Movie Name	Mary's Rating	John's Rating	Do I like it?
Lord of the Rings II	1	5	No
...	...	...	...
Star Wars I	4.5	4	Yes
Gravity	3	3	?

- What am I still missing?
- I need to somehow bind these ratings to my taste.
- Therefore, I am going to label this data as I like it or not.

# Visualized movie data



The question is:  
“Am I going to like Gravity?”



# Write a computer program for predicting it

- Labels:                    “I like it”  $\rightarrow 1$                     “I don’t like it”  $\rightarrow 0$
- Inputs:                    Mary’s rating, John’s rating
- A decision function can be as simple as weighted linear combination of my friends:

$$h_{\theta,b} = \theta_1 x_1 + \theta_2 x_2 + b$$

$$h_{\theta,b} = \theta^T x + b$$

- This function has a problem. Its values are unbounded. We want its output to be in the range of 0 and 1.

# Bound its values between 0 and 1

- Below function is unbounded:


$$h_{\theta,b} = \theta^T x + b$$

- We are going to bound its output:

$$h_{\theta,b} = g(\theta^T x + b),$$

where  $g(z)$  is sigmoid function.

$$g(z) = \frac{1}{1 + \exp(-z)}$$



# Using past data to learn the decision function

- We will use the past data to learn  $\theta, b$  to approximate  $y$ . In particular, we want to obtain  $\theta, b$  such that:

$h_{\theta,b}(x^{(1)}) \approx y^{(1)}$  where  $x^{(1)}$  is my friend's ratings for 1<sup>st</sup> movie.

$h_{\theta,b}(x^{(2)}) \approx y^{(2)}$  where  $x^{(2)}$  is my friend's ratings for 2<sup>nd</sup> movie.

...

$h_{\theta,b}(x^{(m)}) \approx y^{(m)}$  where  $x^{(m)}$  is my friend's ratings for m<sup>th</sup> movie.

# Using past data to learn the decision function

To find values of  $\theta$  and  $b$  we can minimize the following *cost function*:

$$J(\theta, b) = (h_{\theta, b}(x^{(1)}) - y^{(1)})^2 + (h_{\theta, b}(x^{(2)}) - y^{(2)})^2 + \dots + (h_{\theta, b}(x^{(m)}) - y^{(m)})^2$$

$$J(\theta, b) = \sum_{i=1}^m (h_{\theta, b}(x^{(i)}) - y^{(i)})^2$$

Use Stochastic Gradient Descent (SGD):

$$\theta_1 = \theta_1 - \alpha \Delta \theta_1$$

$$\theta_2 = \theta_2 - \alpha \Delta \theta_2$$

$$b = b - \alpha \Delta b$$

# Apply our magic Stochastic Gradient Descent

1. Initialize the parameters  $\theta$  and  $b$  at random
2. Pick a random example  $\{x^{(i)}, y^{(i)}\}$
3. Compute the partial derivatives of  $\theta_1, \theta_2, b$
4. Update parameters using:


$$\theta_1 = \theta_1 - \alpha \Delta \theta_1$$

$$\theta_2 = \theta_2 - \alpha \Delta \theta_2$$

$$b = b - \alpha \Delta b$$

Stop it when parameters don't change much, or after a certain number of iterations.


# Here is my decision boundary




Gravity movie is slightly on the “don’t watch” side.

**With this data set**, it seems like “not watching it” makes more sense.

Another way of representing our model



# Were we lucky? Let's plot Susan's movies.



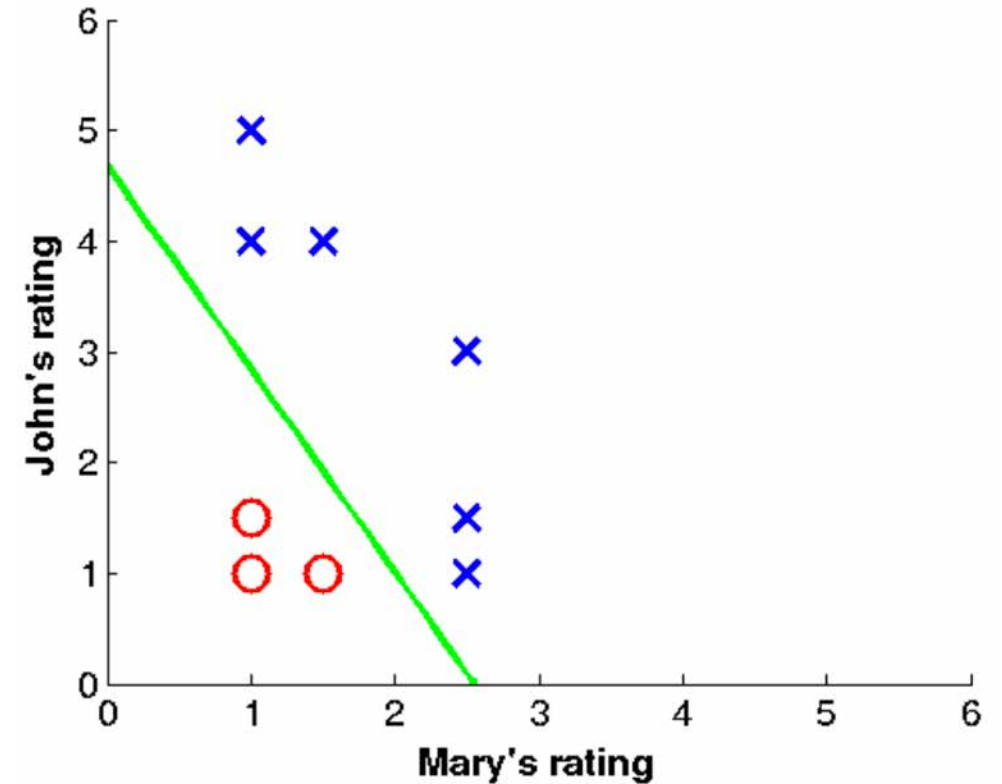

Susan likes some of the movies both Mary and John rated poorly.

How can I have a linear decision boundary separate these?

Maybe we should split the problem into two.




Divide and conquer: We have 2 decision functions.




# A decision function of decision functions

Movie Name	Output by decision function $h_1$	Output by decision function $h_2$	Does Susan like it?
Lord of the Rings II	$h_1(x^{(1)})$	$h_2(x^{(2)})$	No
...	...	...	...
Star Wars I	$h_1(x^{(n)})$	$h_2(x^{(n)})$	Yes
Gravity	$h_1(x^{(n+1)})$	$h_2(x^{(n+1)})$	?


This problem is the same problem as Mary and John case.  
Just consider Output by decision function  $h_1$  and  $h_2$  as two new friends.



A neural network it is!




# A deeper neural network




# Deep vs Shallow Networks

- When the problem does exhibit nonlinear properties, deep networks seem computationally more attractive than shallow networks.
- Bolded edges mean computation paths that need the red neuron to produce the final output.




Deep Network




Shallow Network

# A better activation function: Rectified Linear Unit (ReLU)

$$g(z) = \max(0, z)$$




The sigmoid activation function




The rectified linear activation function

Why is ReLU better?  
Consider derivative of sigmoid.




# Derivative of ReLU





# An even better activation function?

## Parametric Rectified Linear Unit (PReLU)



ReLU vs PReLU

# Softmax Regression

- Softmax regression (or multinomial logistic regression) is a generalization of logistic regression to the case where we want to handle multiple classes.
- In logistic regression we assumed that the labels were binary:
$$y^{(i)} \in \{0, 1\}$$
- Softmax regression allows us to handle
$$y^{(i)} \in \{1, \dots, K\}$$
- where  $K$  is the number of classes.  $K$

# Softmax Regression

- Recall that we have a training set  $\{(x^{(i)}, y^{(i)}), \dots, (x^{(m)}, y^{(m)})\}$  of  $m$  labeled examples.
- Input features are  $x^{(i)} \in \mathbb{R}^n$
- Output features are  $y^{(i)} \in \{1, \dots, K\}$
- For example, in the M-NIST digit recognition task, we would have  $K=10$  different classes.
- Given a test input  $x$ , we want our hypothesis to estimate the probability that  $P(y = k|x)$  for each value of  $k=1, \dots, K$ .

# Softmax Regression

- Given a test input  $\mathbf{x}$ , we want our hypothesis to estimate the probability that  $P(y = k|\mathbf{x})$  for each value of  $k=1, \dots, K$ .
- we want to estimate the probability of the class label taking on each of the  $K$  different possible values.
- Thus, our hypothesis will output a  $K$ -dimensional vector (whose elements sum to 1) giving us our  $K$  estimated probabilities.


$$h_{\theta,b}(\mathbf{x}) = \begin{bmatrix} P(y = 1|\mathbf{x}; \theta, b) \\ P(y = 2|\mathbf{x}; \theta, b) \\ \vdots \\ P(y = K|\mathbf{x}; \theta, b) \end{bmatrix} = \frac{1}{\sum_{j=1}^K \exp(\theta^{(j)T} \mathbf{x} + b^{(j)})} \begin{bmatrix} \exp(\theta^{(1)T} \mathbf{x} + b^{(1)}) \\ \exp(\theta^{(2)T} \mathbf{x} + b^{(2)}) \\ \vdots \\ \exp(\theta^{(K)T} \mathbf{x} + b^{(K)}) \end{bmatrix}$$

# Softmax Regression

$$h_{\theta,b}(x) = \begin{bmatrix} P(y = 1|x; \theta, b) \\ P(y = 2|x; \theta, b) \\ \vdots \\ P(y = K|x; \theta, b) \end{bmatrix} = \frac{1}{\sum_{j=1}^K \exp(\theta^{(j)T} x + b^{(j)})} \begin{bmatrix} \exp(\theta^{(1)T} x + b^{(1)}) \\ \exp(\theta^{(2)T} x + b^{(2)}) \\ \vdots \\ \exp(\theta^{(K)T} x + b^{(K)}) \end{bmatrix}$$

- Notice that  $\frac{1}{\sum_{j=1}^K \exp(\theta^{(j)T} x)}$  is a normalization constant.
- Distribution sums up to 1. This makes it a probability distribution.

# A better neural network model visualization



# Softmax Layer Visualized

